

# MODERN BIG DATA ARCHITECTURES II

PRICE PREDICTION OF BITCOIN IN REAL TIME

## Content Table:

- 1) Problem Statement
- 2) Big Data Pipeline
- 3) Kafka
- 4) Binance API
- 5) Spark Streaming
- 6) ML using Prophet
- 7) InfluxDB
- 8) Conclusion


# Problem Statement

We are an IT consulting company that helps hedge funds to leverage their portfolio of investments by providing real time analytics, for better decision making, and helping them to acquire a perfect hedge.



In this scenario we are working with Two SIGMA to have a better forecasting model by predicting and obtaining the price of Bitcoin in real time while using the **API** from **Binance** and help them achieve a better portfolio management and performance.

# BIG DATA PIPELINE:



# What is Kafka?

Kafka combines three key capabilities:

- 1) To **publish** (write) and **subscribe to** (read) streams of events, including continuous import/export of your data from other systems.
- 2) To **store** streams of events durably and reliably for as long as you want.
- 3) To **process** streams of events as they occur or retrospectively.

Kafka is primarily used to build real-time streaming data pipelines and applications that adapt to the data streams. It combines messaging, storage, and stream processing to allow storage and analysis of both historical and real-time data.




# Start Kafka - Set Up:

```
ignaciogonzalez@Ignacioso-MacBook-Pro-2 ~ % brew install java
```

```
ignaciogonzalez@Ignacioso-MacBook-Pro-2 ~ % brew install kafka
```

- We need to edit the server properties so that we can run kafka server in local  
vim **/usr/local/etc/kafka/server.properties**
- Uncomment the server settings and update the value from
  - o **listeners=PLAINTEXT://:9092**  
to
  - o **listeners=PLAINTEXT://localhost:9092**



```
>_
```

```
ignaciogonzalez@Ignacioso-MacBook-Pro-2 ~ % zookeeper-server-start /usr/local/etc/kafka/zookeeper.properties &
```

```
ignaciogonzalez@Ignacioso-MacBook-Pro-2 ~ % kafka-server-start /usr/local/etc/kafka/server.properties &
```

# Running the Python Script on terminal:

```
(base) ignaciogonzalez@Ignacios-MacBook-Pro-2 Desktop % python3 BinanceData.py "btcbusd" -t binance_data -b localhost:9092
BTCBUSD,1648522311673,46920.37000000
BTCBUSD,1648522312668,46920.37000000
BTCBUSD,1648522313678,46920.37000000
BTCBUSD,1648522314620,46920.37000000
BTCBUSD,1648522315674,46920.37000000
BTCBUSD,1648522316672,46920.37000000
BTCBUSD,1648522317631,46920.37000000
BTCBUSD,1648522318643,46920.37000000
BTCBUSD,1648522319639,46920.37000000
BTCBUSD,1648522320673,46920.37000000
BTCBUSD,1648522321659,46920.37000000
BTCBUSD,1648522322682,46920.37000000
BTCBUSD,1648522323622,46920.37000000
BTCBUSD,1648522324648,46920.37000000
BTCBUSD,1648522325554,46922.93000000
BTCBUSD,1648522326658,46932.27000000
```



# Data Source - Binance API:

- We have used the **push based mechanism** of reading data from Binance API, where we subscribed to Binance API.
- The data returned by Binance API is in form of the JSON
- We read E (timestamp), s (crypto symbol), o (open price) from this JSON data returned by Binance

```
{
  "stream": "ticker",
  "data": {
    "e": "24hrTicker", // Event type
    "E": 123456789, // Event time
    "s": "ABC_0DX-BNB", // Symbol
    "p": "0.0015", // Price change
    "P": "250.00", // Price change percent
    "w": "0.0018", // Weighted average price
    "x": "0.0009", // Previous day's close price
    "c": "0.0025", // Current day's close price
    "Q": "10", // Close trade's quantity
    "b": "0.0024", // Best bid price
    "B": "10", // Best bid quantity
    "a": "0.0026", // Best ask price
    "A": "100", // Best ask quantity
    "o": "0.0010", // Open price
    "h": "0.0025", // High price
    "l": "0.0010", // Low price
    "v": "10000", // Total traded base asset volume
    "q": "18", // Total traded quote asset volume
    "O": 0, // Statistics open time
    "C": 66400000, // Statistics close time
    "F": "0", // First trade ID
    "L": "18150", // Last trade Id
    "n": 18151 // Total number of trades
  }
}
```



```
$ sudo python3 -m pip install binance-connector
```

```
pip install confluent-kafka
```

# Data Source - Binance Python Script :

```
from binance.websocket.spot.websocket_client import SpotWebsocketClient as WebsocketClient
from confluent_kafka import Producer
import argparse
import socket
```


```
def create_data_to_publish(event_raw):
    line = ""
    ts = event_raw['E']
    symbol = event_raw['s']
    open = event_raw['o']
    line = f'{symbol},{ts},{open}'
    return line
```

```
f'{symbol},{ts},{open}'  
'BTCBUSD,1234567890,42000'
```

```
def binance_callback_decorator(producer, topic):
    def stream_callback(message):

        if message != None:
            line = create_data_to_publish(message)
            print(line)
            if producer != None:
                producer.produce(topic, value=line)
                producer.flush()
            else:
                print("no producer")

    return stream_callback
```



# Kafka Producer :


```
if __name__ == "__main__":
    parser = argparse.ArgumentParser()
    parser.add_argument("symbols", help="Comma-separated list of symbols (ex. btcbusd, ethbusd)")
    parser.add_argument("-b", "--broker",
                        help="server:port of the Kafka broker where messages will be published")
    parser.add_argument("-t", "--topic",
                        help="topic where messages will be published")
    args = parser.parse_args()
```

- So we are taking 3 arguments from the user:
  - **Symbol** -> Crypto name for which we want to pull the data from Binance API.
  - **-b / --broker** -> kafka server ip : port where kafka is running
  - **-t / --topic** -> kafka topic name where we want to publish the data.

```
producer = None
topic = args.topic
if args.broker != None:
    conf = {'bootstrap.servers': args.broker, 'client.id': socket.gethostname(), 'broker.address.family':
'v4'}
    producer = Producer(conf)

ws_client = WebSocketClient()
ws_client.start()
```

```
for symbol_raw in args.symbols.split(","):
    symbol = symbol_raw.strip()
    ws_client.instant_subscribe(stream = f"{symbol}@ticker", callback = binance_callback_decorator(producer, topic))
```



# Spark Streaming - Setup :

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % brew install jupyter
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % brew install apache-spark
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % pip install findspark
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % export SPARK_LOCAL_IP="127.0.0.1"
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % export PYSPARK_DRIVER_PYTHON='jupyter'
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % export PYSPARK_DRIVER_PYTHON_OPTS='notebook --no-browser --port=8889'
```

```
pyspark --packages
```



# Spark Session + Spark Streaming :

```
import findspark  
findspark.init()
```

```
import os  
os.environ['PYSPARK_SUBMIT_ARGS'] = '--packages "org.apache.spark:spark-sql-kafka-0-10_2.12:3.0.3" pyspark-shell'  
  
from pyspark.context import SparkContext  
from pyspark.sql.session import SparkSession  
  
spark = (SparkSession.builder  
    .appName("Binance - Hyper price increase detection")  
    .getOrCreate())
```

```
ohlcvDF = spark  
    .readStream  
    .format("kafka")  
    .option("kafka.bootstrap.servers", "localhost:9092")  
    .option("subscribe", "binance_data")  
    .option("startingOffsets", "earliest")  
    .load()
```

```
DataFrame[key: binary, value: binary, topic: string, partition: int, offset: bigint, ti
```



# Spark Streaming :

```
from pyspark.sql.functions import split, col
from pyspark.sql.types import DoubleType, LongType

#TCBUSD,1648234367780,43974.77000000

ohlcvDF = ohlcvDF
    .selectExpr("CAST(value AS STRING)")
    .select(split("value",'\'').alias("fields"))
    .withColumn("timestamp",col("fields").getItem(1).cast(LongType()))
    .withColumn("symbol",col("fields").getItem(0))
    .withColumn("open",col("fields").getItem(2).cast(DoubleType()))
    .drop("fields")|
```



# Spark Streaming :

```
class InfluxDBWriter:
    def __init__(self):
        self.t = 'FqxfLtwynoPObY_Ug9NYmA51_M4vwC5qqIEXn2pjLhM1adsUQJFhvOM33L8iRHuvOtPxAyG-L-V-3ACAcQUTUA=='
        self.o = 'project' # database -> org
        self.b = 'binance' # table -> bucket
        self.client = InfluxDBClient(url="http://localhost:8086", token=self.t, org=self.o) # client creation
        self.write_api = self.client.write_api() # influx db write api

    def open(self, partition_id, epoch_id): # opens the connection to influx db
        print("Opened %d, %d" % (partition_id, epoch_id))
        return True

    def process(self, row): # writing the data into influx db table
        self.write_api.write(bucket=self.b, record=self._row_to_line_protocol(row))

    def close(self, error): #closing write apis and influxdb client
        self.write_api.__del__()
        self.client.__del__()
        print("Closed with error: %s" % str(error))


    def _row_to_line_protocol(self, row: Row): #converting the data into the format how it is stored in influx db
        # TODO map Row to LineProtocol
        p = Point("mess").tag("coin", row['symbol']).field("value", row['open']).time(row['timestamp'], WritePrecision)
        return p
```

# Spark Streaming + Write Stream API:

```
from pyspark.sql.functions import concat, lit  
  
ohlcvDF.writeStream \  
    .format("console") \  
    .start() \  
    .awaitTermination()
```

-----  
Batch: 0  
-----

timestamp	symbol	open
1647938724128	BTCBUSD	41296.64
1647938725124	BTCBUSD	41292.35
1647938726104	BTCBUSD	41292.35
1647938727078	BTCBUSD	41292.35
1647938728117	BTCBUSD	41292.35
1647938729121	BTCBUSD	41292.35
1647938730132	BTCBUSD	41292.35



# Machine Learning (Prophet) - Setup on terminal:

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % pip install pystan~=2.14
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % pip install fbprophet
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % pip install -U scikit-learn scipy matplotlib
```



- As Anaconda for ML and Python was already installed the requirements were already satisfied.

# What is Prophet:

We are using facebook prophet to forecast the price of the crypto for the next 5-10 mins.

Prophet is a procedure for **forecasting time series data** based on an **additive model where non-linear trends are fit** with yearly, weekly, and daily seasonality, plus holiday effects. It works best with time series that have strong seasonal effects and several seasons of historical data. Prophet is **robust** to missing data and shifts in the trend, and typically handles outliers well.



# Machine Learning - Script:

[influxDb -> (python notebook) 5 minutes -> forecasting for next 5 minutes -> influxDb]

```
from pyspark.sql import Row
from datetime import datetime
from influxdb_client import InfluxDBClient, Point, WritePrecision
from influxdb_client.client.write_api import SYNCHRONOUS
```

```
class InfluxDBWriter:
    def __init__(self):
        self.t = 'C3GlPodiT0dqAH9LH_Q8Z0vtXdkZ84DAkFE0phwEtgR_lbc7TYP4Sz6rTavhbNcLIjT9joY15Lyz4x1kkd30LA=='
        self.o = 'project' # database -> org
        self.b = 'binance' # table -> bucket
        self.client = InfluxDBClient(url="http://localhost:8086", token=self.t, org=self.o) # client creation
        self.write_api = self.client.write_api() # influx db write api
```



# Machine Learning:

```
def read_from_influx(coin):
    tables = query_api.query(f'from(bucket: "{b}") |> range(start: -1h,stop: now()) |> filter(fn: (r) => r._measurement == "coin")')
    resultset = []
    for table in tables:
        for row in table.records:
            resultset.append({'time':row['_time'].strftime(datetime_format), 'value':row['_value'], 'coin':row['coin']})
    df = pd.DataFrame(resultset, columns=["time", "value", "coin"])
    return df
```

	time	value
0	2022-03-25 11:54:10	42956.26
1	2022-03-25 11:54:11	42956.26
2	2022-03-25 11:54:12	42956.26
3	2022-03-25 11:54:13	42956.26
4	2022-03-25 11:54:14	42956.25
...	...	...
16786	2022-03-25 17:54:04	43968.00
16787	2022-03-25 17:54:05	43968.01
16788	2022-03-25 17:54:06	43966.69
16789	2022-03-25 17:54:07	43966.68
16790	2022-03-25 17:54:08	43966.69



# Machine Learning:

```
def write_in_influx(df, column_name):
    for ind in df.index:
        p = Point(predicted_m).tag("coin", df['coin'][ind]).field(column_name, df[column_name][ind]).time(int(df
            write_api.write(bucket=b, record=p)
```

```
coin = "BTCBUSD"
df = read_from_influx(coin)

df_train = df[['time', 'value']]
df_train['ds'] = df_train['time']
df_train['y'] = df_train['value']
df_train.drop(['time', 'value'], axis='columns', inplace=True)
print(df_train)

model = Prophet(daily_seasonality=True, weekly_seasonality=True, yearly_seasonality=True)
model.fit(df_train)
```

# Machine Learning:

```
last_t = df_train['ds'].iloc[-1]

future = list()
for i in range(1, prediction_time_in_secs):
    x = to_datetime(last_t) + timedelta(seconds=i)
    future.append(x)

future = DataFrame(future)
future.columns = ['ds']
future['ds'] = to_datetime(future['ds'])

print(future['ds'])

forecast = model.predict(future)
```

```
0  2022-03-25 17:54:09
1  2022-03-25 17:54:10
2  2022-03-25 17:54:11
3  2022-03-25 17:54:12
4  2022-03-25 17:54:13
...
294 2022-03-25 17:59:03
295 2022-03-25 17:59:04
296 2022-03-25 17:59:05
297 2022-03-25 17:59:06
298 2022-03-25 17:59:07
```

# Machine Learning:

```
last_t = df_train['ds'].iloc[-1]

future = list()
for i in range(1, prediction_time_in_secs):
    x = to_datetime(last_t) + timedelta(seconds=i)
    future.append(x)

future = DataFrame(future)
future.columns = ['ds']
future['ds'] = to_datetime(future['ds'])

print(future['ds'])

forecast = model.predict(future)
```

```
0  2022-03-25 17:54:09
1  2022-03-25 17:54:10
2  2022-03-25 17:54:11
3  2022-03-25 17:54:12
4  2022-03-25 17:54:13
...
294 2022-03-25 17:59:03
295 2022-03-25 17:59:04
296 2022-03-25 17:59:05
297 2022-03-25 17:59:06
298 2022-03-25 17:59:07
```

# Machine Learning:

```
future[ 'yhat' ] = forecast[ 'yhat' ]
future[ 'yhat_lower' ] = forecast[ 'yhat_lower' ]
future[ 'yhat_upper' ] = forecast[ 'yhat_upper' ]
future[ 'coin' ] = coin

print(future)

write_in_influx(future, 'yhat')
write_in_influx(future, 'yhat_lower')
write_in_influx(future, 'yhat_upper')
client.close()

time.sleep(sleep_time_in_secs)
```

## What is InfluxDB - Setup :

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % Brew install influxdb
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % brew update
```

```
ignaciogonzalez@Ignacios-MacBook-Pro-2 ~ % pip install influxdb-client  
influxd
```

- We can validate influx db by opening  
<http://localhost:8086>



# What is InfluxDB:

- It is a time-series database (TSDB) optimized for timestamped or time-series data. Time series data are simply measurements or events that are tracked, monitored, downsampled, and aggregated over time.
- InfluxDB is the open-source time-series database.



# Price Prediction of Bitcoin in real time Two Sigma



Add Cell Add Note Variables Annotations

Enable Auto Refresh UTC Past 15m

Data

Explore

Books

Boards

Tasks

Alerts

Settings

Graph



2022-03-29 06:50:00

2022-03-29 06:55:00


2022-03-29 07:00:00

## Conclusion:

In a world that is facing a lot of challenges, being able to provide accurate predictions about the future during times that demand more clarity is key.

In addition, thanks to the democratization of technology, and following the idea of “Open Source” which promotes information age mindset toward innovation.

As a result, the creation of networks of information will keep disrupting the way we live, allowing people around the world to achieve more than ever before.



# THANK YOU RAUL & JORGE

